

An Agentic Retrieval-Augmented Forecasting Framework for India's Automotive Market: Integrating Multi-Document Intelligence, Policy Analysis, and Time-Series Modeling

A PROJECT REPORT

Submitted by

RAVIKANTH AMAJALA [Reg No: RA2412044015067]

Under the Guidance of

Dr. K. VENKATESH

Professor, Department of Networking and Communication

In partial fulfilment of the requirements for the degree of

MASTER OF TECHNOLOGY

in

ARTIFICIAL INTELLIGENCE AND DATA SCIENCE



DEPARTMENT OF COMPUTATIONAL INTELLIGENCE

COLLEGE OF ENGINEERING AND TECHNOLOGY

SRM INSTITUTE OF SCIENCE AND TECHNOLOGY

KATTANKULATHUR – 603 203

MAY 2026

**SRM INSTITUTE OF SCIENCE AND TECHNOLOGY****KATTANKULATHUR –603 203****BONAFIDE CERTIFICATE**

Certified that this M.Tech project report titled “**An Agentic Retrieval-Augmented Forecasting Framework for India’s Automotive Market: Integrating Multi-Document Intelligence, Policy Analysis, and Time Series Modeling**” is the bonafide work of **Ms. Ravikanth Amajala [Reg. No.: RA2412044015067]** who carried out the project work under my supervision. Certified further, that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or award was conferred on an earlier occasion for this or any other candidate.

SIGNATURE**SUPERVISOR****Dr. K. VENKATESH**

Professor

Department of Networking and Communication

SIGNATURE**HEAD OF THE DEPARTMENT****Dr. R. ANNIE UTHRA**

Professor

Department of Computational Intelligence

INTERNAL EXAMINER**EXTERNAL EXAMINER**



DEPARTMENT OF COMPUTATIONAL INTELLIGENCE

SRM Institute of Science and Technology Own Work Declaration Form

Degree/ Course : M.Tech in Artificial Intelligence and Data Science

Student Names : Ravikanth Amajala

Registration Number : RA2412044015067

Title of Work : An Agentic Retrieval-Augmented Forecasting Framework for India's Automotive Market: Integrating Multi-Document Intelligence, Policy Analysis, and Time-Series Modeling

I hereby certify that this assessment compiles with the University's Rules and Regulations relating to Academic misconduct and plagiarism, as listed in the University Website, Regulations, and the Education Committee guidelines.

I confirm that all the work contained in this assessment is my own except where indicated, and that I have met the following conditions:

- Clearly references / listed all sources as appropriate
- Referenced and put in inverted commas all quoted text (from books, web, etc.)
- Given the sources of all pictures, data etc. that are not my own
- Not made any use of the report(s) or essay(s) of any other student(s) either past or present
- Acknowledged in appropriate places any help that I have received from others (e.g. fellow students, technicians, statisticians, external sources)
- Compiled with any other plagiarism criteria specified in the Course handbook / University website

I understand that any false claim for this work will be penalized in accordance with the University policies and regulations.

DECLARATION:

I am aware of and understand the University's policy on Academic misconduct and plagiarism and I certify that this assessment is my own work, except were indicated by referring, and that I have followed the good academic practices noted above.

Ravikanth Amajala RA2412044015067

ACKNOWLEDGEMENT

I express my humble gratitude to **Dr. C. Muthamizhchelvan**, Vice-Chancellor, SRM Institute of Science and Technology, for the facilities extended for the project work and his continued support. I extend my sincere thanks to Dean-CET, SRM Institute of Science and Technology, **Dr. M. Leenus J Martin**, for his invaluable support.

I wish to thank **Dr. Revathi Venkataraman**, Professor & Chairperson, School of Computing, SRM Institute of Science and Technology, for her support throughout the project work. I sincerely appreciate **Dr. R. Annie Uthra**, Professor and Head, Professor, Department of Computational Intelligence, SRM Institute of Science and Technology, for her guidance and support that has been invaluable to my success. I am incredibly grateful to my Program Coordinator, **Dr. S. Ganesh Kumar**, Professor, Department of Data Science and Business Systems, SRM Institute of Science and Technology, for his suggestions and encouragement at all the stages of the project work.

I want to convey my thanks to my Project Coordinator, **Dr. Pitchaimanickam**, Assistant Professor and the Panel Head, **Dr. K. Venkatesh**, Professor, Department of Networking and Communications, SRM Institute of Science and Technology, for their inputs during the project reviews and support.

I register my immeasurable thanks to my Program Manager, **Mr. Samuel Godshine**, Program Manager, Great Learning for leading and helping us to complete the course.

My inexpressible respect and thanks to the guide, **Dr. K. Venkatesh**, Professor, Department of Networking and Communication, SRM Institute of Science and Technology, for providing me with an opportunity to pursue my project under his mentorship. He provided us with the freedom and support to explore the research topics of interest. His passion for solving problems and making a difference in the world has always been inspiring.

I sincerely thank the Department of Computational Intelligence staff and students, SRM Institute of Science and Technology, for their help during my project. Finally, I would like to thank parents, family members, and friends for their unconditional love, constant support, and encouragement.

Ravikanth Amajala [RA2412044015067]

TABLE OF CONTENTS

<i>C.NO.</i>	<i>TITLE</i>	<i>PAGE NO.</i>
	Abstract	vii
	List of Tables	viii
	List of Figures	viii
	List of Symbols and Abbreviations	ix
1.	INTRODUCTION	11
1.1	General	11
1.2	Purpose of the Project	12
1.3	Scope of the Study	13
1.4	Overview of Artificial Intelligence	13
1.5	Machine Learning and Time-Series Forecasting	17
1.6	Retrieval-Augmented Generation (RAG)	18
1.7	Agent-Based Systems	20
2	LITERATURE REVIEW	22
2.1	Traditional Forecasting Models (ARIMA, SARIMA)	22
2.2	Machine Learning in Forecasting	24
2.3	Deep Learning Approaches (LSTM, RNN)	26
2.4	Retrieval-Augmented Systems	28
2.5	Multi-Document Intelligence	31
2.6	Research Gap Analysis	33
3	PROPOSED METHODOLOGY	34
3.1	Dataset Description	34
3.2	Data Preprocessing	35
3.3	Feature Engineering	37
3.4	ARIMA Model	39
3.5	Boosting Models	41
3.6	Retrieval Module (RAG)	44
3.7	Agent-Based Architecture	47
3.8	Algorithm	52
3.9	Justification of Proposed Approach	53

4	RESULTS AND ANALYSIS	54
	4.1 Performance Metrics (MAE, RMSE, MAPE)	54
	4.2 Comparative Analysis	56
	4.3 Graphical Analysis	59
	4.4 Performance Improvement After Optimization	64
5	CONCLUSION	65
	5.1 Limitations	66
6	FUTURE SCOPE	67
7	REFERENCES	69
	APPENDIX 1	70
	APPENDIX 2	72
	PAPER PUBLICATION STATUS	75
	PLAGIARISM REPORT	76

ABSTRACT

Accurate forecasting of the automotive market is critical for strategic planning in the presence of dynamic policy changes, economic fluctuations, and evolving consumer behavior. Traditional time-series forecasting models primarily rely on structured historical data and often fail to incorporate unstructured information such as policy documents and industry reports. To address this limitation, this paper proposes an **Agentic Retrieval-Augmented Forecasting Framework** for predicting trends in India's automotive market over a long-term horizon.

The proposed approach integrates **time-series forecasting models** with **Retrieval-Augmented Generation (RAG)** to combine structured and unstructured data sources. A hybrid modeling strategy employing **AutoRegressive Integrated Moving Average (ARIMA)** and machine learning boosting techniques is utilized to enhance predictive performance. Additionally, an **agent-based architecture** is introduced to coordinate data preprocessing, document retrieval, and forecasting processes, enabling contextual understanding and adaptive decision-making.

The framework leverages multi-document intelligence by extracting relevant insights from policy documents, market analyses, and economic indicators. Performance is evaluated using standard metrics, including **Mean Absolute Error (MAE)**, **Root Mean Square Error (RMSE)**, and **Mean Absolute Percentage Error (MAPE)**. Experimental results demonstrate that the proposed hybrid model outperforms conventional approaches in terms of accuracy and robustness.

The proposed system provides a scalable and intelligent solution for context-aware forecasting and can assist policymakers, industry stakeholders, and researchers in making informed decisions. Furthermore, it establishes a foundation for future research in integrating artificial intelligence with domain-specific knowledge for enhanced predictive analytics.

LIST OF TABLES

2.1	Parameters of ARIMA Model.....	23
2.2	Comparison of Traditional and Machine Learning Models.....	25
2.4	Summary of Retrieval-Augmented Systems.....	30
3.1	Dataset Description and Features.....	35
3.2	Data Preprocessing Techniques.....	37
3.4	ARIMA Model Parameters and Configuration.....	41
3.5	Boosting Model Parameters.....	43
3.6	RAG Retrieval Components.....	45
4.1	Performance Metrics Comparison (MAE, RMSE, MAPE).....	55
4.2	Model Performance Before and After Optimization.....	57
4.4	Performance Improvement After Optimization.....	64

LIST OF FIGURES

1.4	Types of Machine Learning	
	(a) Supervised Learning	14
	(b) Unsupervised Learning	15
	(c) Reinforced Learning.....	16
1.5	Overview of Time-Series Forecasting Process	18
1.6	Architecture of Retrieval-Augmented Generation (RAG)	19
1.7	Agent-Based System Architecture	21
2.1	Flow Diagram of ARIMA Model	24
2.3	Structure of LSTM Network	27
2.4	Working of Retrieval-Augmented Systems	29
2.5	Multi-Document Intelligence Pipeline	32
3.2	Data Preprocessing Workflow	36
3.3	Feature Engineering Process	38
3.4	ARIMA Model Workflow	40
3.5	Boosting Model Architecture	42
3.6	RAG Retrieval Pipeline	45
3.7	(a)Proposed Agentic Framework Architecture	48
	(b)Interaction between Agents and Modules	48
	(c)User Interface for Smart Document Upload and Forecasting.....	49

	(d)Forecast Input Interface.....	50
	(e)Processing Status and Extraction Results Interface.....	51
	(f)Policy Simulation Module Interface.....	51
3.8	Algorithm Flowchart of Proposed System	53
4.1	Predicted vs Actual Values Graph	56
4.2	Error Comparison Graph (MAE, RMSE, MAPE)	58
4.3	(a)Trend Analysis of Automotive Market	60
	(b)Performance Improvement After Optimization	61
	(c)Visualization of Forecast Output Dashboard	62
	(d)Policy Scenario Simulation and Forecast Analysis.....	63
	(e)Policy Simulation Results Summary and Strategic Insights.....	63

LIST OF SYMBOLS AND ABBREVIATIONS

AI	Artificial Intelligence
ML	Machine Learning
DL	Deep Learning
RAG	Retrieval-Augmented Generation
ARIMA	AutoRegressive Integrated Moving Average
SARIMA	Seasonal AutoRegressive Integrated Moving Average
LSTM	Long Short-Term Memory
RNN	Recurrent Neural Network
XGBoost	Extreme Gradient Boosting
GBM	Gradient Boosting Machine
MAE	Mean Absolute Error
RMSE	Root Mean Square Error
MAPE	Mean Absolute Percentage Error
NLP	Natural Language Processing
LLM	Large Language Model
API	Application Programming Interface
GDP	Gross Domestic Product

EV	Electric Vehicle
CSV	Comma-Separated Values
JSON	JavaScript Object Notation
CPU	Central Processing Unit
GPU	Graphics Processing Unit
RAM	Random Access Memory
P	Autoregressive Order (ARIMA)
D	Differencing Order (ARIMA)
Q	Moving Average Order (ARIMA)
$\hat{y}$	Predicted Value
y	Actual Value
ε	Error Term
i.e.	id est
e.g.	exempli gratia
Eq	Equation

CHAPTER 1

INTRODUCTION

This chapter introduces the fundamental concepts of machine learning and time-series forecasting, which form the core of the proposed system. It provides an overview of how data-driven approaches can be used to analyze historical trends and predict future outcomes. The chapter also discusses the motivation behind integrating traditional statistical models with modern machine learning techniques and highlights the importance of forecasting in the automotive domain. Additionally, it outlines the objectives, scope, and structure of the project.

1.1. General

Indian automotive industry is considered as one of the largest, and fastest growing sectors as it has contributed to the economic development, creation of employment and industrial growth of the country. The demand of vehicles has been on a steadily growing trend throughout the years with the increasing urbanization, rising disposable incomes and the swift technological advancements. Also, the shift to electric mobility, the introduction of more strict emission standards, and the development of new policies of the state have brought new dynamics to the automotive ecosystem.

With the ever increasing complexity of the automotive market, there is a rising need to properly forecast demand and trends. The classical forecasting techniques are mainly based on past numerical data like sales data, and economic indicators. Although these methods are useful in giving a baseline knowledge, they cannot always reflect the impact of the external forces, including government laws, policy alterations, industry analysis, and market moods. This weakness leads to a less accurate forecast, particularly in the long term forecasting and strategic planning.

Over the last few years, new possibilities in the improvement of forecasting systems have been introduced by the accessibility of large amounts of unstructured information such as policy documents, research reports and market analyses. But one of the main challenges is to successfully combine this information with structured data. The rise of artificial intelligence, machine learning, and natural language processing methods offers an avenue to help solve this problem by allowing systems to process and make sense of structured and unstructured data sources.

This project is driven by the desire to have an intelligent and scalable forecasting system which breaks the shackles of the conventional methods. The solution presented proposes an Agentic Retrieval-Augmented Forecasting Framework, a combination of time-series forecasting models and Retrieval-Augmented Generation (RAG) to gain a contextual knowledge of various documents. The system takes a hybrid approach to modeling the data based on statistical techniques like ARIMA and machine learning techniques to enhance the accuracy of the prediction.

Moreover, a system based on an agent architecture is created with the purpose of orchestrating the work of different system parts such as data preprocessing, document retrieval and forecasting modules. This facilitates dynamic decision-making and increases the flexibility of the system to evolving market conditions. The proposed framework will offer more reliable and context-aware predictions by incorporating multi-source data and using sophisticated methods of analysis.

The results of this project should enlighten the policymakers, automotive manufacturers and investors to make informed decisions. It also adds to the larger body of smart forecasting systems showing that

artificial intelligence and domain-specific knowledge can be effectively combined.

1.2. Purpose of the Project

The primary purpose of this project is to develop an intelligent and efficient forecasting system capable of predicting trends in India's automotive market by integrating both structured and unstructured data sources. With the increasing complexity of the automotive sector, driven by rapid technological advancements, policy changes, and evolving consumer preferences, there is a growing need for a system that can provide accurate and context-aware predictions.

Traditional forecasting methods are limited in their ability to incorporate external influences such as government regulations, industry reports, and economic indicators. These methods often rely solely on historical numerical data, which restricts their effectiveness in capturing real-world dynamics. The purpose of this project is to overcome these limitations by introducing a hybrid framework that combines time-series forecasting models with advanced artificial intelligence techniques.

This project aims to utilize **Retrieval-Augmented Generation (RAG)** to extract relevant insights from multiple documents, including policy papers and market analyses, and integrate them into the forecasting process. By doing so, the system enhances its understanding of contextual factors that influence market behavior. Additionally, the use of **agent-based architecture** enables efficient coordination between different modules, ensuring seamless data processing and improved adaptability.

Another key objective of the project is to improve prediction accuracy through the use of hybrid modeling techniques such as **ARIMA and machine learning boosting algorithms**. These models are designed to complement each other, capturing both linear and non-linear patterns in the data.

Overall, the purpose of this project is to provide a scalable, intelligent, and cost-effective forecasting solution that can assist stakeholders such as policymakers, automotive companies, and investors in making informed decisions. It also aims to contribute to the advancement of forecasting methodologies by demonstrating the effectiveness of integrating artificial intelligence with domain-specific knowledge.

1.3. Scope of the Study

The scope of this project focuses on the development and evaluation of an intelligent forecasting framework for analyzing trends in the Indian automotive market using advanced artificial intelligence techniques. The study primarily aims to integrate structured time-series data with unstructured textual information to improve the accuracy and reliability of market predictions.

The project covers the use of historical automotive data, including vehicle sales, economic indicators, and market trends, along with unstructured data sources such as government policies, industry reports, and research publications. By incorporating these diverse data sources, the system is designed to provide a more comprehensive understanding of the factors influencing the automotive sector.

The proposed framework utilizes a hybrid approach that combines statistical models, such as **ARIMA**, with machine learning techniques, including boosting algorithms, to capture both linear and non-linear patterns in the data. Additionally, the integration of **Retrieval-Augmented Generation (RAG)** enables the system to extract relevant contextual information from multiple documents, enhancing the overall forecasting capability.

The scope of the study also includes the design of an **agent-based architecture**, which coordinates various components such as data preprocessing, document retrieval, and forecasting modules. This ensures efficient data flow, scalability, and adaptability of the system to changing market conditions.

However, the project is limited to the analysis of available datasets and documents, and its performance depends on the quality and relevance of the input data. Real-time data integration and deployment in a live production environment are beyond the current scope but can be considered for future enhancements.

Overall, this study provides a foundation for developing intelligent, context-aware forecasting systems and demonstrates their application in the automotive domain.

1.4 Overview of Artificial Intelligence

Artificial Intelligence (AI) is a branch of computer science that focuses on the development of systems capable of performing tasks that typically require human intelligence. These tasks include learning from data, reasoning, problem-solving, decision-making, and understanding natural language. AI has emerged as a transformative technology, influencing a wide range of domains such as healthcare, finance, transportation, and manufacturing.

At its core, AI aims to create intelligent systems that can analyze large volumes of data, identify patterns, and make predictions or decisions with minimal human intervention. With the rapid growth of data availability and computational power, AI techniques have become increasingly effective in solving complex real-world problems.

AI can be broadly categorized into three types: **Narrow AI**, **General AI**, and **Super AI**. Narrow AI refers to systems designed to perform specific tasks, such as recommendation systems or speech recognition. General AI represents a more advanced concept where machines possess the ability to perform any intellectual task that a human can do, although such systems are still under development. Super AI is a theoretical concept where machines surpass human intelligence in all aspects.

A significant subset of AI is **Machine Learning (ML)**, which enables systems to learn from data and improve their performance over time without being explicitly programmed. ML algorithms can be classified into supervised learning, unsupervised learning, and reinforcement learning, depending on the type of data and learning approach. Another important subset is **Deep Learning (DL)**, which uses artificial neural networks with multiple layers to model complex patterns and relationships in data.

In the context of this project, AI plays a crucial role in enabling intelligent forecasting by integrating structured numerical data with unstructured textual information. Techniques such as natural language processing (NLP) allow the system to extract meaningful insights from documents, while machine learning models enhance prediction accuracy. The combination of these approaches enables the development of a context-aware forecasting system that adapts to changing market conditions.

Overall, Artificial Intelligence provides the foundation for building advanced, data-driven systems capable of solving complex problems and supporting informed decision-making across various domains.




Figure 1.4(a) supervised learning

Figure 1.4(a) This figure illustrates the supervised learning process where labeled data is used to train a model. The algorithm learns the mapping between inputs and outputs, enabling accurate predictions on unseen data.




Figure 1.4.(b) unsupervised learning

Figure 1.4(b) This figure represents unsupervised learning, where the model analyzes unlabeled data to identify hidden patterns and structures. It is commonly used for clustering and data grouping tasks.




Figure 1.4.(c) reinforcement learning

Figure 1.4(c) This figure shows the reinforcement learning framework, where an agent interacts with an environment and improves its decisions through rewards and penalties over time.

1.5 Machine Learning and Time-Series Forecasting

Machine Learning (ML) is a subset of Artificial Intelligence that enables systems to learn patterns from data and make predictions or decisions without being explicitly programmed. It has become a fundamental tool for analyzing large datasets and solving complex problems across various domains. In the context of forecasting, machine learning techniques provide the ability to model both linear and non-linear relationships, thereby improving prediction accuracy compared to traditional statistical methods.

Time-series forecasting is a technique used to predict future values based on previously observed data points collected over time. In the automotive industry, time-series data may include vehicle sales, demand trends, fuel prices, and economic indicators. Accurate forecasting of such data is essential for strategic planning, inventory management, and policy formulation.

Traditional time-series models, such as the **AutoRegressive Integrated Moving Average (ARIMA)** model, are widely used due to their effectiveness in capturing linear trends and seasonality in data. ARIMA models rely on past values and error terms to make predictions, making them suitable for structured datasets with consistent patterns. However, these models often struggle to capture complex non-linear relationships and external influences.

Machine learning approaches, including **Decision Trees, Random Forests, and Gradient Boosting algorithms**, have been introduced to address these limitations. These models can handle high-dimensional data and capture intricate patterns that traditional models may overlook. Additionally, advanced techniques such as **Recurrent Neural Networks (RNN)** and **Long Short-Term Memory (LSTM)** networks are specifically designed for sequential data and are capable of learning temporal dependencies over long periods.

In this project, a hybrid approach is adopted by combining traditional time-series models with machine learning techniques. The ARIMA model is used as a baseline for capturing linear patterns, while boosting algorithms enhance the model's ability to learn non-linear relationships. This combination improves overall forecasting performance and robustness.

Furthermore, integrating machine learning with external data sources, such as policy documents and market reports, enhances the contextual understanding of the forecasting system. This enables more accurate and reliable predictions, especially in dynamic environments where multiple factors influence market behavior.

Overall, machine learning and time-series forecasting together form a powerful framework for predictive analytics, enabling the development of intelligent systems that support data-driven decision-making.




Figure 1.5 Overview of Time-Series Forecasting Process

Figure 1.5. This figure presents the overall workflow of time-series forecasting, including data collection, preprocessing, model training, and prediction. It highlights how temporal patterns are utilized for future estimation.

1.6 Retrieval-Augmented Generation (RAG)

Retrieval-Augmented Generation (RAG) is an advanced artificial intelligence framework that combines information retrieval techniques with generative models to enhance the quality, accuracy, and contextual relevance of outputs. Unlike traditional models that rely solely on pre-trained knowledge, RAG systems dynamically retrieve relevant information from external data sources and incorporate it into the generation process.

The core idea behind RAG is to bridge the gap between structured and unstructured data by enabling

systems to access and utilize large collections of documents such as policy reports, research papers, and industry analyses. This is particularly important in domains where up-to-date and context-specific information plays a critical role in decision-making.

A typical RAG system consists of two main components: a **retriever** and a **generator**. The retriever is responsible for searching and identifying relevant documents or passages from a knowledge base using techniques such as semantic search and vector embeddings. The generator then uses this retrieved information, along with the input query, to produce more accurate and context-aware outputs.

In the context of this project, RAG is used to enhance forecasting by incorporating insights from multiple external sources, including government policies, market reports, and economic analyses. These documents often contain valuable information that influences automotive market trends but is not captured in structured numerical data. By integrating this information, the forecasting system gains a deeper understanding of the underlying factors affecting the market.

The use of RAG also improves the interpretability of the system by providing contextual explanations for predictions. This is particularly useful for stakeholders who require not only accurate forecasts but also an understanding of the reasons behind them. Furthermore, RAG enables the system to remain adaptable, as it can continuously incorporate new information without requiring complete retraining.

Overall, Retrieval-Augmented Generation plays a crucial role in developing intelligent, context-aware forecasting systems. By combining retrieval mechanisms with predictive models, it enhances both the accuracy and relevance of predictions, making it a key component of the proposed framework.




Figure 1.6: Architecture of Retrieval-Augmented Generation (RAG)

Figure 1.6: This figure depicts the RAG architecture, where relevant information is retrieved from a document database and combined with a generative model to produce context-aware outputs.

1.7 Agent-Based Systems

Agent-Based Systems are a class of computational systems in which multiple autonomous entities, known as agents, interact with each other and their environment to perform complex tasks. Each agent is designed to operate independently, make decisions based on available information, and collaborate with other agents to achieve a common objective. This paradigm is widely used in distributed systems, artificial intelligence, and complex problem-solving environments.

An agent typically possesses key characteristics such as **autonomy, adaptability, reactivity, and goal-oriented behavior**. Autonomy allows the agent to function without continuous human intervention, while adaptability enables it to adjust its actions based on changing conditions. Reactivity refers to the ability of the agent to respond to environmental inputs in real time, and goal-oriented behavior ensures that the agent works towards achieving specific objectives.

In modern intelligent systems, agent-based architectures are used to divide complex workflows into smaller, manageable components. Each agent is responsible for a specific task, such as data collection, preprocessing, analysis, or decision-making. This modular approach improves system scalability, flexibility, and efficiency, as individual components can be updated or modified without affecting the entire system.

In the context of this project, an **agent-based architecture** is employed to coordinate various components of the forecasting framework. Different agents are assigned specific roles, such as handling data preprocessing, performing document retrieval using RAG, executing forecasting models, and generating outputs. These agents interact with each other to ensure smooth data flow and efficient execution of tasks.

The use of agent-based systems enhances the adaptability of the framework, allowing it to respond dynamically to new data and changing market conditions. It also improves fault tolerance, as the failure of one agent does not necessarily disrupt the entire system. Furthermore, this architecture supports scalability, making it suitable for handling large datasets and integrating additional functionalities in the future.

Overall, Agent-Based Systems provide a robust and flexible foundation for building intelligent and distributed applications. Their integration into the proposed framework enables efficient coordination of multiple processes, resulting in improved performance and more reliable forecasting outcomes.




Figure 1.7: Agent-Based System Architecture

Figure 1.7 This figure illustrates the architecture of an agent-based system, where multiple intelligent agents interact with different system components to perform specific tasks. Each agent is responsible for functions such as data processing, retrieval, forecasting, and decision-making. The collaborative interaction between agents enables efficient workflow management and improves the overall adaptability and performance of the system.

CHAPTER 2

LITERATURE REVIEW

In this section, we discuss the previous research that has been conducted in the field of automotive market forecasting and predictive analytics. The information presented in this section includes advancements in time-series forecasting models, machine learning techniques, and hybrid intelligent systems that aim to improve prediction accuracy. It also covers the integration of structured and unstructured data sources, along with recent developments in Retrieval-Augmented Generation (RAG) and agent-based frameworks.

Several studies have focused on improving forecasting performance by combining statistical models with machine learning approaches, as well as incorporating external factors such as economic indicators, policy changes, and market dynamics. These approaches highlight the importance of contextual awareness in predictive systems. While traditional models provide a strong mathematical foundation, recent research emphasizes the need for intelligent systems capable of adapting to complex and dynamic environments.

Two important aspects are often considered in forecasting studies—**prediction accuracy** and **model robustness**. Prediction accuracy refers to how closely the predicted values match the actual observations, while robustness indicates the model's ability to maintain performance under varying conditions and datasets. These factors are typically evaluated using metrics such as MAE, RMSE, and MAPE.

2.1. Traditional Forecasting Models (ARIMA, SARIMA)

Traditional forecasting models have been widely used for analyzing time-series data due to their strong statistical foundation and interpretability. Among these, the **AutoRegressive Integrated Moving Average (ARIMA)** model is one of the most commonly used techniques for forecasting linear trends in time-series data.

The ARIMA model is defined by three parameters: **p (autoregressive order)**, **d (differencing order)**, and **q (moving average order)**. It works by capturing relationships between past observations and

error terms to predict future values. ARIMA is particularly effective for stationary time-series data and is widely applied in domains such as economics, finance, and sales forecasting.

An extension of ARIMA, known as **Seasonal ARIMA (SARIMA)**, is used to model seasonal patterns present in time-series data. SARIMA incorporates additional seasonal parameters to account for repeating patterns over fixed intervals, making it suitable for datasets with periodic fluctuations.

Despite their effectiveness, these models have certain limitations. They primarily focus on linear relationships and may not perform well when the data exhibits complex non-linear patterns. Additionally, traditional models do not incorporate external factors such as policy changes, market sentiment, or unstructured textual information, which can significantly influence real-world outcomes.

Due to these limitations, recent research has shifted towards hybrid approaches that combine traditional statistical models with machine learning techniques. These approaches aim to enhance forecasting accuracy by leveraging the strengths of both methodologies.

Table 2.1: Parameters of ARIMA Model

Parameter	Name	Description	Role in Model
p	AutoRegressive (AR) Order	Number of lag observations included in the model	Captures relationship between current value and past values
d	Differencing Order	Number of times the data is differenced to make it stationary	Removes trend and stabilizes mean
q	Moving Average (MA) Order	Number of lagged forecast errors in the model	Captures noise and short-term fluctuations
AR(p)	Autoregressive Component	Uses past values of the series	Models linear dependency over time
I(d)	Integrated Component	Applies differencing to data	Ensures stationarity
MA(q)	Moving Average Component	Uses past forecast errors	Smoothens random variations

Table 2.1. describes the key parameters of the ARIMA model, including autoregressive, differencing, and moving average components. These parameters collectively help in modeling time-series data by capturing temporal dependencies, removing non-stationarity, and accounting for noise.




Figure 2.1: Flow Diagram of ARIMA Model

Figure 2.1. shows the workflow of the ARIMA model, including stationarity checking, differencing, parameter selection, model training, and forecasting. It demonstrates how ARIMA captures temporal dependencies in time-series data.

2.2. Machine Learning in Forecasting

Machine Learning (ML) techniques have gained significant attention in recent years for their ability to improve forecasting accuracy by modeling complex and non-linear relationships in data. Unlike traditional statistical models, which rely on predefined assumptions, machine learning algorithms learn patterns directly from data, making them highly adaptable to diverse and dynamic environments.

In the context of time-series forecasting, machine learning models are capable of handling large volumes of data with multiple features, including economic indicators, market trends, and external influencing factors. Algorithms such as **Decision Trees, Random Forests, Support Vector Machines (SVM), and Gradient Boosting methods** have been widely used for predictive analysis. These models can capture intricate relationships between variables and provide more accurate predictions compared to traditional approaches in many real-world scenarios.

One of the key advantages of machine learning models is their ability to perform **feature engineering and selection**, which helps in identifying the most relevant variables influencing the prediction. This improves both the efficiency and interpretability of the forecasting system. Additionally, ensemble methods such as **Random Forest and Gradient Boosting** combine multiple models to enhance prediction performance and reduce overfitting.

Despite their advantages, machine learning models also face certain challenges. They often require large amounts of high-quality data for training and may lack interpretability compared to traditional statistical models. Furthermore, these models typically do not inherently account for temporal dependencies unless specifically designed for sequential data.

To address these limitations, machine learning techniques are often combined with traditional time-series models or deep learning approaches. Hybrid models leverage the strengths of both methodologies, capturing linear trends using statistical models and non-linear patterns using machine learning algorithms.

In this project, machine learning plays a crucial role in enhancing the forecasting framework by complementing traditional models such as ARIMA. The integration of boosting algorithms improves prediction accuracy and enables the system to adapt to complex patterns in automotive market data. This combination forms a robust and efficient forecasting system capable of handling real-world challenges.

Table 2.2: Comparison of Traditional and Machine Learning Models

Aspect	Traditional Models (ARIMA, SARIMA)	Machine Learning Models (Regression, RF, XGBoost)
Data Type Handling	Univariate or limited multivariate	Highly multivariate and heterogeneous
Linearity Assumption	Strict linear relationships assumed	Non-parametric; no strict assumption
Non-Linearity	Limited; requires data transformation	High capability to model complex patterns
Feature Engineering	Low; based on lags and residuals	High; requires manual feature selection
Interpretability	High; statistical coefficients are clear	Moderate to Low (Black-box nature)
Accuracy	High for small, stationary datasets	High for large, complex datasets
Computational Complexity	Low; computationally inexpensive	High; requires significant resources
Scalability	Limited scalability to Big Data	Excellent scalability with parallel processing

Aspect	Traditional Models (ARIMA, SARIMA)	Machine Learning Models (Regression, RF, XGBoost)
Use Cases	Economic forecasting, stock trends	Fraud detection, image recognition, sales

Table 2.2 presents a comparison between traditional statistical models and machine learning approaches for forecasting. It highlights differences in their ability to handle linear and non-linear patterns, scalability, and adaptability to complex datasets.

2.3. Deep Learning Approaches (LSTM, RNN)

Deep Learning is an advanced subset of machine learning that utilizes artificial neural networks with multiple layers to model complex patterns and relationships in data. It has shown remarkable success in handling large-scale and high-dimensional datasets, particularly in applications involving sequential and time-dependent data.

Among the various deep learning techniques, **Recurrent Neural Networks (RNN)** are specifically designed for processing sequential data. RNNs maintain a form of memory by passing information from one step to the next, allowing them to capture temporal dependencies in time-series data. This makes them suitable for forecasting tasks where past observations influence future outcomes. However, standard RNNs face challenges such as the **vanishing gradient problem**, which limits their ability to learn long-term dependencies.

To overcome this limitation, **Long Short-Term Memory (LSTM)** networks were introduced as an improved version of RNNs. LSTM networks incorporate specialized components called **gates**—namely the input gate, forget gate, and output gate—which regulate the flow of information through the network. These gates enable LSTM models to retain important information over longer sequences and discard irrelevant data, making them highly effective for time-series forecasting.

LSTM models have been widely applied in domains such as financial forecasting, weather prediction, and demand analysis due to their ability to capture both short-term and long-term dependencies. They can model non-linear relationships and complex temporal patterns that traditional statistical models and basic machine learning algorithms may fail to capture.

Despite their advantages, deep learning models require large volumes of data and significant computational resources for training. They are also relatively complex and may lack interpretability compared to simpler models. Additionally, deep learning approaches primarily rely on structured numerical data and may not directly incorporate unstructured information such as textual documents or policy reports.

In the context of this project, deep learning approaches provide valuable insights into advanced forecasting techniques. While the primary framework focuses on hybrid models combining ARIMA and machine learning, the understanding of RNN and LSTM architectures highlights the potential for future enhancements. Integrating deep learning models with external knowledge sources, such as through Retrieval-Augmented Generation (RAG), can further improve the capability of forecasting systems.

Overall, deep learning approaches, particularly RNN and LSTM, represent a powerful class of models for time-series forecasting, offering the ability to capture complex temporal dynamics and improve predictive performance in data-driven applications.




Figure 2.3: Structure of LSTM Network

Figure 2.3 This figure illustrates the internal structure of a Long Short-Term Memory (LSTM) network, highlighting its key components such as the input gate, forget gate, output gate, and memory cell. These elements work together to control the flow of information, enabling the model to retain important long-term dependencies and overcome the vanishing gradient problem in sequential data processing.

2.4.Retrieval-Augmented Systems

Retrieval-Augmented Systems represent a significant advancement in artificial intelligence by combining information retrieval techniques with learning-based models to enhance performance and contextual understanding. These systems are designed to overcome the limitations of traditional models that rely solely on pre-trained knowledge by dynamically accessing external data sources during execution.

The core concept of retrieval-augmented systems involves two primary components: a **retrieval module** and a **processing or generation module**. The retrieval module is responsible for identifying and extracting relevant information from a large corpus of documents using techniques such as keyword-based search, semantic similarity, and vector embeddings. The retrieved information is then passed to the processing module, which utilizes it to improve the quality and relevance of predictions or generated outputs.

Recent research has shown that retrieval-augmented approaches significantly improve performance in tasks that require domain-specific knowledge and contextual awareness. By incorporating external sources such as research papers, policy documents, and industry reports, these systems can adapt to new information without requiring complete retraining. This makes them particularly useful in dynamic environments where knowledge is continuously evolving.

In the context of forecasting, retrieval-augmented systems enable the integration of **unstructured textual data** with structured numerical datasets. For example, policy changes, economic reports, and market analyses often contain valuable insights that influence trends but are not captured in traditional datasets. By retrieving and incorporating this information, the forecasting model gains a more comprehensive understanding of the factors affecting predictions.

One of the key advantages of retrieval-augmented systems is their ability to provide **explainable outputs**. Since the system retrieves relevant documents to support its predictions, it becomes easier to interpret and justify the results. This is particularly important in decision-making scenarios where transparency and reliability are essential.

However, these systems also face challenges such as the need for efficient indexing, retrieval accuracy, and handling large-scale document collections. The performance of the system depends heavily on the quality of the retrieval process and the relevance of the selected documents.

In this project, retrieval-augmented systems form the foundation of the proposed framework by enabling the integration of multi-document intelligence into the forecasting process. This approach enhances prediction accuracy, improves contextual understanding, and provides a scalable solution for real-world applications.




Figure 2.4: Working of Retrieval-Augmented Systems

Figure 2.4. This figure illustrates how retrieval-augmented systems fetch relevant documents based on query similarity and use them to enhance the quality of generated responses.

Table 2.4: Summary of Retrieval-Augmented Systems

Component	Description	Technique Used	Purpose
User Query	Initial natural language input from the end user	Natural Language Processing	Define the information need for the system
Query Encoding	Conversion of raw text into numerical format	Vector representation	Enable mathematical comparison with stored data
Embedding Model	Neural network used to generate dense vectors	Transformer-based embeddings	Capture semantic meaning within high-dimensional space
Vector Database	Storage system for indexed document embeddings	Efficient indexing architectures	Allow rapid retrieval from large external datasets
Similarity Search	Identifying relevant data points near the query vector	Cosine similarity	Measure relevance based on semantic proximity
Document Retrieval	Extraction of top-ranked text chunks from the corpus	Semantic search	Provide the generator with factual evidence
Context Aggregation	Integration of retrieved facts with the original query	Prompt augmentation	Construct a comprehensive input for the model
Generator (LLM)	Large Language Model producing the final response	In-context learning	Synthesize response using contextual knowledge

Table 2.4. summarizes the key components and functionalities of retrieval-augmented systems. It outlines how external knowledge sources are integrated with machine learning models to enhance contextual understanding and improve output relevance.

2.5. Multi-Document Intelligence

Multi-Document Intelligence refers to the ability of a system to extract, analyze, and integrate information from multiple heterogeneous documents to generate meaningful insights. In many real-world applications, valuable information is distributed across various sources such as research papers, policy documents, industry reports, and news articles. Individually, these documents provide partial knowledge; however, when combined, they offer a comprehensive understanding of a given domain.

Traditional data-driven systems primarily rely on structured datasets and often overlook the rich information available in unstructured textual sources. Multi-document intelligence addresses this limitation by leveraging techniques from **natural language processing (NLP)** and information retrieval to process large volumes of textual data. This enables systems to identify relevant patterns, relationships, and contextual signals that are not present in numerical data alone.

One of the key challenges in multi-document intelligence is the efficient retrieval and integration of relevant information from diverse sources. Documents may vary in format, structure, and quality, making it difficult to extract consistent and accurate insights. Techniques such as **text embedding, semantic similarity, and document ranking** are commonly used to identify the most relevant information for a given task.

In the context of forecasting, multi-document intelligence plays a crucial role in enhancing prediction accuracy by incorporating external knowledge. For instance, government policies, economic reports, and market analyses significantly influence automotive industry trends. By analyzing multiple documents, the system can capture these external factors and integrate them into the forecasting process, leading to more informed and context-aware predictions.

Furthermore, multi-document intelligence improves the robustness and adaptability of the system. As new documents become available, the system can continuously update its knowledge base without requiring extensive retraining. This ensures that the forecasting model remains relevant in a rapidly changing environment.

In this project, multi-document intelligence is implemented through the integration of Retrieval-Augmented Generation (RAG), which enables the system to retrieve and utilize information from multiple sources effectively. This approach enhances the overall capability of the forecasting framework by combining structured data analysis with contextual knowledge derived from

unstructured documents.

Overall, multi-document intelligence provides a powerful mechanism for leveraging diverse information sources, improving decision-making, and enabling the development of intelligent, knowledge-driven systems.




Figure 2.5: Multi-Document Intelligence Pipeline

Figure 2.5 This figure presents a pipeline where multiple documents are processed, indexed, and analysed to extract meaningful insights. It enables integration of diverse data sources for better decision-making.

2.6. Research Gap Analysis

Existing forecasting systems primarily rely on structured numerical data such as historical sales and economic indicators. While these approaches are effective for capturing trends, they exhibit several limitations:

- Dependence only on structured datasets
- Lack of integration with unstructured data such as policy documents and reports
- Limited ability to adapt to real-world policy changes
- Absence of scenario-based simulation for decision-making

To address these limitations, the proposed system introduces a hybrid and intelligent framework with the following improvements:

- Integration of structured (time-series) and unstructured (policy documents) data
- Use of Retrieval-Augmented Generation (RAG) for contextual understanding
- Implementation of an agent-based architecture for modular processing
- Support for policy-driven scenario simulation and forecasting

This combination enables a more robust, adaptive, and context-aware forecasting system compared to traditional approaches.

CHAPTER 3

PROPOSED

METHODOLOGY

In this section, we provide the details of the dataset used, the proposed architecture for the forecasting framework, and the techniques employed to predict trends in the automotive market. The system is designed to integrate both structured and unstructured data sources to improve prediction accuracy and contextual understanding.

First, a dataset consisting of historical automotive sales data, economic indicators, and related structured information is collected and processed. In addition to this, a collection of unstructured documents such as government policies, industry reports, and market analyses is incorporated into the system. Secondly, using a combination of statistical models and machine learning techniques, a hybrid forecasting model is developed to capture both linear and non-linear patterns in the data. Furthermore, a Retrieval-Augmented Generation (RAG) module is used to extract relevant contextual information from multiple documents.

Finally, an agent-based architecture is employed to coordinate the different components of the system, including data preprocessing, document retrieval, forecasting, and result generation. This integrated approach enables the system to produce accurate and context-aware predictions for the automotive market.

3.1.Dataset Description

The dataset used in this project consists of both structured and unstructured data sources relevant to the automotive market. The structured data includes historical records such as vehicle sales, production volumes, fuel types, and economic indicators like inflation rates and gross domestic product (GDP). These datasets are essential for capturing numerical trends and patterns over time.

In addition to structured data, unstructured data sources such as government policy documents, industry reports, and market research publications are also utilized. These documents contain valuable contextual information that influences market behavior, such as regulatory changes, technological advancements, and shifts in consumer preferences.

The structured dataset is typically stored in formats such as CSV or Excel files, while unstructured data is collected in the form of text documents and PDFs. Before being used in the forecasting process, the data undergoes preprocessing steps including data cleaning, handling missing values, normalization, and transformation into suitable formats.

The dataset is divided into training and testing sets to evaluate the performance of the forecasting models. The training set is used to build and optimize the models, while the testing set is used to validate their predictive accuracy.

By combining both structured and unstructured data, the dataset provides a comprehensive foundation

for building an intelligent forecasting system capable of capturing real-world dynamics in the automotive market

Table 3.1 – Dataset Description and Features

Feature Name	Description	Data Type	Unit	Role in Model
Date	Temporal marker for monthly observations	Date	YYYY-MM-DD	Input Feature
Vehicle Sales	Total monthly units sold across dealerships	Integer	Units	Target Variable
Fuel Price	Average retail price of gasoline/diesel	Float	₹ per Litre	Input Feature
Economic Index	Consumer confidence or purchasing power metric	Float	Index Value	Input Feature
Policy Impact	Binary indicator for active tax or subsidy changes	Boolean	0/1	Input Feature
Sales Growth	Percentage change from previous period	Float	%	Derived Feature
Interest Rate	Average auto loan rates for the period	Float	%	Input Feature

Table 3.1 provides an overview of the dataset used in this study, including key features and variables. It highlights the structure of the data and the attributes considered for forecasting and analysis.

3.2.Data Preprocessing

Data preprocessing is a crucial step in the development of any forecasting system, as the quality and consistency of the input data directly influence the performance of the models. In this project, preprocessing is applied to both structured and unstructured data to ensure that the information is clean, relevant, and suitable for analysis.

For structured data, preprocessing involves several steps such as **data cleaning, handling missing values, normalization, and transformation**. Inconsistent or incomplete records are identified and either corrected or removed to maintain data integrity. Missing values are handled using appropriate techniques such as mean imputation or interpolation, depending on the nature of the dataset. Additionally, normalization and scaling techniques are applied to ensure that all features are on a

comparable scale, which is particularly important for machine learning models.

Time-series data requires special preprocessing to ensure stationarity and consistency. Techniques such as **differencing, smoothing, and trend removal** are applied to stabilize the data and make it suitable for models like ARIMA. Seasonal patterns, if present, are also identified and handled appropriately.

For unstructured data, preprocessing involves the use of **natural language processing (NLP)** techniques. This includes text cleaning, tokenization, removal of stop words, and conversion of text into numerical representations using methods such as word embeddings. These steps enable the system to extract meaningful information from documents such as policy reports and market analyses.

In the context of Retrieval-Augmented Generation (RAG), preprocessing also includes the creation of **document embeddings and indexing**, which allows efficient retrieval of relevant information during the forecasting process. Documents are processed and stored in a vector database, enabling semantic search based on similarity.

Overall, data preprocessing ensures that both structured and unstructured data are transformed into a consistent and usable format. This step plays a vital role in improving model accuracy, reducing noise, and enabling effective integration of multiple data sources within the proposed forecasting framework.




Figure 3.2: Data Preprocessing Workflow

Figure 3.2. This figure shows the preprocessing steps applied to raw data, including cleaning, transformation, normalization, and feature preparation. These steps ensure data quality for model training.

Table 3.2: Data Preprocessing Techniques

Technique	Description	Method Used	Purpose
Data Cleaning	Removal of noise and inconsistent data points	Deduplication / Noise filtering	Improve data quality and model reliability
Missing Value Handling	Estimating values for gaps in time-series data	Linear Interpolation / Mean Imputation	Maintain temporal continuity for forecasting
Outlier Detection	Identification of extreme statistical deviations	Z-score / Interquartile Range (IQR)	Prevent skewed training and biased results
Data Transformation	Modifying distribution to reach stationarity	Log Transformation / Differencing	Normalize variance and stabilize mean
Normalization / Scaling	Adjusting feature ranges to a common scale	Min-Max Scaling / Standardization	Ensure uniform feature contribution to gradients
Feature Encoding	Converting categorical variables to numerical	One-Hot Encoding / Label Encoding	Enable mathematical operations on non-numeric data
Data Integration	Consolidating data from multiple sources	Schema Mapping / Key-based Merging	Provide a holistic view for multivariate analysis

Table 3.2 outlines the preprocessing techniques applied to the dataset, including cleaning, normalization, and transformation methods. These steps ensure data quality and improve model performance.

3.3.Feature Engineering

Feature engineering is a critical step in the forecasting process, as it involves the extraction, selection, and transformation of relevant variables that significantly influence model performance. The objective of this step is to convert raw data into meaningful features that can effectively represent underlying patterns and relationships in the dataset.

In this project, feature engineering is applied to both structured time-series data and unstructured textual data. For structured data, features are derived from historical automotive records such as vehicle sales, production volumes, fuel types, and economic indicators. Additional features such as **lag variables, moving averages, and trend indicators** are generated to capture temporal dependencies and improve the predictive capability of the models.

Time-based features play an important role in forecasting tasks. Attributes such as **year, month, and**

seasonal indicators are extracted to identify recurring patterns and trends in the data. These features help the model understand periodic fluctuations and long-term growth patterns in the automotive market.

For machine learning models, feature selection techniques are used to identify the most relevant variables that contribute to prediction accuracy. This helps in reducing dimensionality, improving computational efficiency, and minimizing the risk of overfitting. Techniques such as correlation analysis and importance ranking are employed to select optimal features.

In addition to structured data, feature engineering is also performed on unstructured textual data using natural language processing techniques. Relevant information from policy documents and market reports is transformed into numerical representations through methods such as **text embeddings and semantic vectorization**. These features capture contextual information that cannot be obtained from numerical data alone.

The integration of features from both structured and unstructured sources enables the system to develop a comprehensive understanding of the factors influencing automotive market trends. This hybrid feature set enhances the ability of the forecasting models to capture both quantitative patterns and qualitative insights.

Overall, feature engineering plays a vital role in improving model performance by ensuring that the input data is informative, relevant, and well-structured for the forecasting process.




Figure 3.3: Feature Engineering Process

Figure 3.3: This figure illustrates the feature engineering process applied to raw data, including extraction, transformation, and creation of relevant features such as lag variables, rolling averages, and encoded attributes. These engineered features enhance the model’s ability to capture patterns and improve forecasting accuracy.

3.4. ARIMA Model

The AutoRegressive Integrated Moving Average (ARIMA) model is a widely used statistical technique for time-series forecasting. It is particularly effective for modeling linear relationships in data and capturing temporal dependencies based on past observations. In this project, ARIMA is used as a baseline model to analyze historical trends in the automotive dataset.

The ARIMA model is defined by three key parameters: p , d , and q , where p represents the order of the autoregressive component, d denotes the degree of differencing required to make the series stationary, and q indicates the order of the moving average component. The model can be represented as $\text{ARIMA}(p, d, q)$.

The autoregressive (AR) component models the relationship between the current value and its previous values, while the moving average (MA) component captures the relationship between the current value and past error terms. The integrated (I) component involves differencing the data to remove trends and ensure stationarity.

Mathematically, the ARIMA model can be expressed as:

$$Y_t = c + \sum_{i=1}^p \phi_i Y_{t-i} + \sum_{i=1}^q \theta_i \epsilon_{t-i} + \epsilon_t$$

where:

Y_t is the value at time t

c is a constant term

ϕ_i are autoregressive coefficients

θ_i are moving average coefficients

ϵ_t is the error term

Before applying the ARIMA model, the time-series data is tested for stationarity using methods such as the Augmented Dickey-Fuller (ADF) test. If the data is non-stationary, differencing is applied to stabilize the mean and variance.

The model parameters (p , d , q) are selected using techniques such as Autocorrelation Function (ACF) and Partial Autocorrelation Function (PACF) plots. These tools help identify the appropriate lag values for the autoregressive and moving average components.

Although ARIMA is effective for capturing linear trends, it has limitations in handling non-linear

patterns and external influences. Therefore, in this project, ARIMA is combined with machine learning techniques to form a hybrid model that improves overall forecasting performance.

Overall, the ARIMA model serves as a fundamental component of the proposed framework, providing a strong statistical baseline for time-series prediction.

Table 3.4 presents the configuration of the ARIMA model used in this study, including selected parameter values. These configurations are essential for achieving optimal forecasting performance.




Figure 3.4: ARIMA Model Workflow

Figure 3.4. This figure illustrates the detailed ARIMA workflow, covering visualization, stationarity check, differencing, parameter tuning, and forecast generation.

Table 3.4 – ARIMA Model Parameters and Configuration

Parameter	Description	Value Used	Selection Method	Purpose
p (AR Order)	Number of lag observations in the model	2	PACF Analysis / AIC	Account for relationship between current and past values
d (Differencing)	Degree of differencing for stationarity	1	ADF Test / Visual Inspection	Remove trend and seasonality from the time-series
q (MA Order)	Size of the moving average window	1	ACF Analysis / BIC	Capture relationship between observation and error lag
Stationarity Test	Statistical test for unit root presence	$p < 0.05$	Augmented Dickey-Fuller (ADF)	Validate if the series is mean-stationary
ACF Analysis	Correlation between series and its lags	Lag-1 cut-off	Correlogram Plot	Identify the optimal order of the MA component
PACF Analysis	Partial correlation of series and lags	Lag-2 cut-off	Correlogram Plot	Identify the optimal order of the AR component
Model Selection	Evaluation of model fit vs. complexity	Minimised Score	AIC / BIC Criteria	Select the most parsimonious model configuration

3.5.Boosting Model

Boosting is a powerful ensemble machine learning technique that improves prediction accuracy by combining multiple weak learners to form a strong predictive model. Unlike individual models, boosting methods focus on reducing errors by iteratively learning from previous mistakes, making them highly effective for capturing complex and non-linear patterns in data.

In this project, boosting models are used to enhance the forecasting performance by complementing the limitations of traditional statistical approaches such as ARIMA. While ARIMA effectively captures linear trends in time-series data, boosting algorithms are capable of modeling non-linear relationships and interactions between multiple features.

One of the commonly used boosting techniques is **Gradient Boosting**, which builds models sequentially by minimizing a loss function. Each new model is trained to correct the errors made by the previous models. Another widely used variant is **Extreme Gradient Boosting (XGBoost)**, which offers improved efficiency, regularization, and scalability, making it suitable for large datasets and real-

world applications.

The boosting process involves assigning higher importance to data points that are difficult to predict. During each iteration, the model adjusts its parameters to reduce the overall prediction error. This iterative refinement results in a highly accurate and robust model.

In the context of this project, boosting models are trained using features derived from both structured data (such as sales and economic indicators) and engineered features (such as lag values and trends). These models help capture complex dependencies that cannot be effectively modeled by traditional time-series techniques alone.

One of the key advantages of boosting models is their ability to handle high-dimensional data and perform implicit feature selection. They also provide flexibility in modeling different types of data distributions. However, boosting methods may require careful tuning of hyperparameters and can be computationally intensive.

By integrating boosting models with ARIMA in a hybrid framework, the system benefits from both statistical and machine learning approaches. This combination improves prediction accuracy, enhances robustness, and enables the forecasting system to adapt to complex real-world scenarios.

Overall, boosting models play a crucial role in the proposed methodology by strengthening the predictive capability of the system and addressing the limitations of traditional forecasting techniques.




Figure 3.5: Boosting Model Architecture

Figure 3.5: This figure illustrates the architecture of the boosting model, where multiple weak learners are trained sequentially to improve overall prediction performance. Each subsequent model focuses on correcting the errors of the previous ones, resulting in a strong ensemble model capable of capturing complex patterns and reducing prediction error.

Table 3.5 – Boosting Model Parameters

Parameter	Description	Value Used	Purpose
Number of Estimators	Total sequential boosting stages (trees) to be performed	100	Controls model complexity and capacity to learn from residuals
Learning Rate	Step size shrinkage used in update to prevent overfitting	0.1	Weights the contribution of each tree via gradient descent
Maximum Depth	Maximum number of levels in each individual tree	5	Limits the depth to capture interactions without over-fitting
Subsample Ratio	Fraction of training instances used for fitting each tree	0.8	Introduces stochasticity to improve generalization and speed
Loss Function	Objective function to be optimized during training	Squared Error	Measures the difference between predicted and actual values
Regularization (Alpha)	L1 regularization term applied to leaf weights	0.01	Encourages sparsity and reduces model sensitivity to outliers
Min Samples Split	Minimum samples required to split an internal node	2	Acts as a stopping criterion to control tree granularity

Table 3.5 lists the key parameters of the boosting model and their configurations. These parameters influence the model’s ability to capture complex patterns and improve predictive accuracy.

The hybrid model combines ARIMA and Boosting predictions using weighted averaging. This approach improves performance by capturing both linear time dependencies and complex non-linear interactions in the data.

The ARIMA model is effective in capturing linear trends and seasonality in time-series data but lacks the capability to model complex non-linear relationships. On the other hand, Boosting models such as XGBoost are powerful in learning non-linear patterns but do not explicitly account for temporal dependencies.

To overcome these limitations, a hybrid approach is adopted, combining ARIMA and Boosting predictions using weighted averaging. This allows the model to leverage the strengths of both approaches—temporal pattern recognition from ARIMA and non-linear learning from Boosting—resulting in improved forecasting accuracy and robustness.

3.6 Retrieval Module (RAG)

The Retrieval-Augmented Generation (RAG) module plays a crucial role in enhancing the forecasting framework by integrating external knowledge with predictive modeling. Unlike traditional models that rely solely on historical data, the RAG approach incorporates relevant information from multiple documents to improve contextual understanding and prediction accuracy.

The retrieval module is designed to fetch the most relevant information from a structured vector database based on a given query. Initially, the input query is transformed into a vector representation using an embedding model. This allows semantic comparison rather than simple keyword matching. The generated query vector is then used to perform similarity search within the vector database, which contains preprocessed and embedded documents such as policy reports, market analyses, and economic indicators.

Once the similarity search is performed, the system retrieves the top relevant documents based on similarity scores. These documents are then filtered and ranked to ensure only the most meaningful and contextually relevant information is retained. The selected documents are aggregated to form a structured context that can be effectively utilized by downstream components.

The primary advantage of the RAG-based retrieval module lies in its ability to incorporate dynamic, real-world knowledge into the forecasting process. This enables the system to account for external factors such as policy changes, economic conditions, and market trends, which are not always captured in historical numerical data.

Furthermore, the retrieval process ensures scalability and adaptability, as new documents can be continuously added to the database without requiring retraining of the entire model. This makes the system robust and suitable for real-world applications where data is constantly evolving.

In summary, the RAG retrieval module enhances the overall system by providing context-aware insights, improving prediction quality, and enabling the integration of multi-document intelligence into the forecasting framework.




Figure 3.6: RAG Retrieval Pipeline

Figure 3.6. This figure shows the retrieval pipeline where user queries are converted into embeddings, matched with a vector database, and relevant documents are retrieved for context generation.

Table 3.6: RAG Retrieval Components

Component	Description	Technique / Model Used	Role in System
User Query	Natural language request or question from the end user	Textual String Input	Define the initial semantic search intent
Query Encoding	Transformation of text into high-dimensional numerical data	Vector Representation	Convert text into a machine-readable format

Component	Description	Technique / Model Used	Role in System
Embedding Model	Deep learning model used to generate dense embeddings	Sentence Transformers (e.g., BERT)	Capture semantic meaning and relationships
Vector Database	Efficient storage and indexing for large-scale embeddings	FAISS / Pinecone	Enable rapid retrieval from external datasets
Similarity Search	Finding top-K vectors nearest to the query vector	Cosine Similarity / Euclidean Distance	Identify the most relevant information chunks
Document Retrieval	Extraction of raw text associated with matched vectors	Semantic Search	Retrieve factual source content for generation
Relevance Filtering	Post-processing to remove noise or irrelevant results	Threshold Filtering / Re-ranking	Optimize the precision of retrieved context
Context Preparation	Formatting data for the Language Model prompt	Contextual Knowledge Integration	Provide grounding for high-quality generation

Table 3.6 describes the components involved in the RAG retrieval process, including embedding models, vector storage, and retrieval mechanisms. These components enable efficient integration of external knowledge into the system.

In this implementation, cosine similarity is used for document retrieval due to its simplicity and efficiency. However, in real-world deployments, scalable vector databases such as FAISS or Chroma can be utilized to handle large-scale document collections and improve retrieval performance.

In this implementation, cosine similarity is used for document retrieval due to its simplicity and effectiveness. However, for large-scale real-world applications, scalable vector databases such as FAISS or Chroma can be integrated to improve retrieval efficiency, speed, and scalability.

3.7 Agent-Based Architecture

The proposed system adopts an agent-based architecture to efficiently coordinate multiple components involved in forecasting and retrieval. This approach divides the overall system into specialized agents, each responsible for a specific task, thereby improving modularity, scalability, and flexibility.

At the core of the architecture lies the Coordination Agent, which acts as the central controller. It manages the flow of data and ensures proper interaction between different agents. Upon receiving input data or user queries, the Coordination Agent determines the appropriate processing path and delegates tasks to the respective agents.

The Data Preprocessing Agent is responsible for preparing raw data for analysis. It performs operations such as data cleaning, handling missing values, normalization, and transformation. This ensures that the data is consistent and suitable for further processing.

The Feature Engineering Agent extracts meaningful features from the preprocessed data. It generates derived attributes such as lag features, moving averages, and trend indicators, which enhance the performance of forecasting models.

The Retrieval Agent implements the Retrieval-Augmented Generation (RAG) mechanism. It converts queries into vector representations and performs similarity search on a vector database to retrieve relevant documents. These documents provide contextual knowledge that complements the numerical forecasting models.

The Forecasting Agent applies time-series and machine learning models such as ARIMA and boosting techniques to generate predictions. It uses both engineered features and contextual insights obtained from the retrieval module to improve forecasting accuracy.

Finally, the outputs from different agents are combined and processed by the Output Module, which presents the final predictions and insights in a structured and interpretable format.

The agent-based architecture offers several advantages, including parallel processing, ease of system updates, and improved maintainability. Each agent can be modified or enhanced independently without affecting the entire system. This makes the architecture highly suitable for real-world applications where data sources and requirements continuously evolve.

In summary, the agent-based design enables efficient integration of data processing, retrieval, and forecasting, resulting in a robust and intelligent system capable of delivering accurate and context-aware predictions.




Figure 3.7.(a) – Proposed Agentic Framework Architecture

Figure 3.7 This figure presents the overall system architecture, integrating data processing, retrieval modules, forecasting models, and decision-making components into a unified framework.




Figure 3.7.(b): Interaction between Agents and Modules

Figure 3.7.(b) This figure illustrates the interaction between different system components, highlighting the flow of information between agents, data processing modules, and forecasting units.




Figure 3.7.(c) User Interface for Smart Document Upload and Forecasting

Figure 3.7.(c) Illustrates the user interface of the proposed system, which enables users to upload multiple document formats such as PDF, Excel, and CSV files. The interface follows a step-based workflow, guiding users through document upload, forecasting, and result visualization. It also provides backend status monitoring and supports seamless integration with the agentic retrieval and forecasting modules.

 **Step 2: What Should We Forecast?**

 **Your documents have been uploaded and indexed in the vector database.**
Now ask what you'd like to forecast from the data, and our AI models will analyze and predict future trends.

 What would you like to forecast?

e.g., EV sales, revenue, units shipped, market growth

 We'll automatically find and analyze the best matching data series from your documents.

 How many months ahead should we forecast?

36 months

 Recommended: 36 months (3 years) for optimal predictions

 **Generate AI-Powered Forecast**

Figure 3.7.(d): Forecast Input Interface

Figure 3.7(d) This interface allows users to specify forecasting requirements after document ingestion. The system leverages indexed data from the vector database and enables users to define forecasting parameters such as target variable and prediction horizon. The input is then processed by the forecasting models to generate future predictions.

 **Processing Complete**

8
Files Processed

0
With Forecasts

0
Models Trained

China NEV-policyupdate-jul2020.pdf
Status: **SUCCESS**

Extraction Details:
 Time Series Found: 1

China-green-future-ev-jan2021.pdf
Status: **FAILED**

Extraction Details:



Figure 3.7.(e): Processing Status and Extraction Results Interface

Figure 3.7(e) This figure presents the system interface displaying the processing status and extraction results. It shows how input documents are analyzed, processed, and converted into structured information, along with real-time status updates. The interface enhances transparency by allowing users to monitor progress and view extracted insights effectively.




Figure 3.7.(f): Policy Simulation Module Interface

Figure 3.7(f) This interface represents the interactive policy simulation module, where users can adjust key policy parameters such as infrastructure investment, subsidies, and manufacturing incentives. The system dynamically evaluates the impact of these factors on EV market growth, enabling scenario-based forecasting and decision support.

3.8 Algorithm

This section presents the step-by-step algorithm followed in the proposed agentic retrieval-augmented forecasting framework. The algorithm integrates data preprocessing, feature engineering, retrieval-based contextual enhancement, and forecasting to generate accurate and context-aware predictions.

Algorithm: Hybrid Forecasting with RAG Integration

Input: Time-series dataset, Policy documents

Output: Forecasted values

1. Perform data preprocessing (missing value handling, normalization)
2. Apply feature engineering (lag features, rolling averages)
3. Encode documents using embedding techniques
4. Perform similarity-based retrieval of relevant documents
5. Train ARIMA model on time-series data
6. Train Boosting model on engineered features
7. Generate predictions from both models
8. Combine predictions using weighted averaging
9. Evaluate model performance using MAE, RMSE, and MAPE
10. Generate final forecast output




Figure 3.8: Algorithm Flowchart of Proposed System

Figure 3.8 This figure illustrates the step-by-step algorithmic flow of the proposed system, starting from data input and preprocessing to model training, retrieval, and forecast generation. It highlights the integration of ARIMA, boosting, and retrieval-augmented components, followed by evaluation using performance metrics. The flowchart provides a clear representation of how the system processes data and produces accurate forecasting outputs.

3.9 Justification of Proposed Approach

The proposed hybrid framework is designed to overcome the limitations of standalone models. ARIMA effectively captures linear trends and seasonality but struggles with non-linear relationships. Boosting models (e.g., XGBoost) capture complex non-linear patterns but may overlook temporal dependencies. By combining both, the hybrid model leverages complementary strengths—ARIMA captures linear temporal patterns, while boosting models capture non-linear relationships—resulting in improved accuracy and robustness.

Additionally, integrating Retrieval-Augmented Generation (RAG) enables the system to incorporate unstructured knowledge (e.g., policy documents), which traditional forecasting models typically ignore. This improves contextual awareness and decision support.

CHAPTER 4

RESULTS

This chapter presents the experimental results and performance evaluation of the proposed forecasting system. It includes graphical visualizations, error metrics, and comparative analysis of different models such as ARIMA, boosting, and the hybrid approach. The results are interpreted to highlight improvements in prediction accuracy and model efficiency. Furthermore, the impact of policy scenarios and simulation outputs is analyzed to demonstrate the practical applicability of the system.

4.1 Performance Metrics (MAE, RMSE, MAPE)

The performance of the proposed forecasting framework is evaluated using standard error metrics that measure the accuracy of predictions. These metrics provide a quantitative assessment of how closely the predicted values match the actual observed values. In this project, three widely used evaluation metrics are considered: **Mean Absolute Error (MAE)**, **Root Mean Square Error (RMSE)**, and **Mean Absolute Percentage Error (MAPE)**.

Mean Absolute Error (MAE) measures the average magnitude of errors between predicted and actual values without considering their direction. It is calculated as:

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

where y_i represents the actual value and $\hat{y}_i$ represents the predicted value. A lower MAE indicates better model performance.

Root Mean Square Error (RMSE) measures the square root of the average of squared differences between predicted and actual values. It gives higher weight to larger errors, making it sensitive to outliers. It is defined as:

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

RMSE is useful for understanding the variance in prediction errors and is commonly used for model comparison.

Mean Absolute Percentage Error (MAPE) expresses the prediction error as a percentage of the actual values, providing a normalized measure of accuracy. It is calculated as:

$$MAPE = \frac{100}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right|$$

MAPE is particularly useful for interpreting the accuracy of the model in relative terms.

These metrics are used to evaluate the performance of different models, including the ARIMA model, boosting model, and the proposed hybrid framework. By comparing the values of MAE, RMSE, and MAPE, the effectiveness of each model can be assessed.

Lower values across these

metrics indicate higher prediction accuracy and better model performance.

The experimental results demonstrate that the proposed hybrid model, which integrates ARIMA, boosting techniques, and retrieval-augmented intelligence, achieves lower error values compared to traditional approaches. This indicates that the combination of structured data analysis and contextual information significantly improves forecasting accuracy.

Overall, the use of these performance metrics provides a comprehensive evaluation of the proposed system and validates its effectiveness in predicting automotive market trends.

Table 4.1: Performance Metrics Comparison (MAE, RMSE, MAPE)

Model	MAE	RMSE	MAPE (%)	Interpretation
ARIMA (Baseline)	425.50	512.30	12.45	Baseline accuracy
XGBoost Regressor	285.20	340.15	8.12	Moderate accuracy
Random Forest	310.45	375.80	9.35	Moderate accuracy
Proposed Hybrid (RAG + Forecasting)	142.30	185.60	4.20	Superior accuracy

Table 4.1 compares the performance of different models using evaluation metrics such as MAE, RMSE, and MAPE. It highlights the effectiveness of each model in terms of prediction accuracy.




Figure 4.1: Predicted vs Actual Values Graph

Figure 4.1: This figure compares the predicted values generated by the model with the actual observed values over time. The close alignment between the two curves indicates the accuracy and reliability of the forecasting model in capturing underlying data patterns.

4.2 Comparative Analysis

A comparative analysis is conducted to evaluate the performance of the proposed hybrid forecasting framework against traditional and machine learning-based models. The models considered for comparison include the **ARIMA model**, **boosting models**, and the **proposed Agentic Retrieval-Augmented hybrid model**.

The ARIMA model serves as a baseline for comparison, as it is widely used for time-series forecasting and effectively captures linear trends in the data. However, it is limited in handling non-linear patterns and does not incorporate external contextual information. As a result, its performance is comparatively lower when applied to complex real-world datasets.

Boosting models, such as Gradient Boosting and XGBoost, demonstrate improved

performance over ARIMA by capturing non-linear relationships and interactions between multiple features. These models are capable of handling high-dimensional data and provide better prediction accuracy. However, they primarily rely on structured data and do not utilize unstructured information such as policy documents and market reports.

The proposed hybrid model combines the strengths of both ARIMA and boosting techniques while integrating **Retrieval-Augmented Generation (RAG)** to incorporate contextual knowledge from multiple documents. This approach enables the system to consider both numerical trends and external influencing factors, resulting in more accurate and reliable predictions.

The comparison of models based on performance metrics such as **MAE, RMSE, and MAPE** indicates that the proposed hybrid model achieves the lowest error values among all approaches. This demonstrates its superior ability to capture complex patterns and adapt to dynamic market conditions.

Furthermore, the hybrid model provides enhanced interpretability by linking predictions with relevant contextual information retrieved from external sources. This makes the system more useful for decision-making compared to standalone models.

Overall, the comparative analysis highlights that while traditional and machine learning models offer valuable insights, the integration of retrieval-based intelligence and agent-based architecture significantly improves forecasting performance. The proposed framework outperforms existing approaches in terms of accuracy, robustness, and contextual understanding, making it a more effective solution for automotive market forecasting.

Table 4.2: Model Performance Before and After Optimization

Model	Metric	Before Optimization	After Optimization	Improvement (%)
ARIMA	MAE	450.2	425.5	5.49
ARIMA	RMSE	545.8	512.3	6.14
ARIMA	MAPE (%)	13.2	12.45	5.68
XGBoost Regressor	MAE	345.8	285.2	17.52
XGBoost Regressor	RMSE	410.5	340.15	17.14
XGBoost Regressor	MAPE (%)	9.8	8.12	17.14
Proposed Hybrid Model	MAE	185.4	142.3	23.25

Proposed Hybrid Model	RMSE	242.1	185.6	23.34
Proposed Hybrid Model	MAPE (%)	5.6	4.2	25

Table 4.2 presents a comparison of model performance before and after optimization. The results demonstrate the improvement achieved through tuning and hybrid modeling techniques.




Figure 4.2: Error Comparison Graph (MAE, RMSE, MAPE)

Figure 4.2 This figure presents a comparison of error metrics—MAE, RMSE, and MAPE—across different models. It highlights the effectiveness of the proposed hybrid model in reducing prediction errors compared to ARIMA and boosting models, demonstrating improved forecasting accuracy.

The results indicate that the Hybrid model consistently outperforms individual models across all evaluation metrics. Compared to ARIMA, the Hybrid model reduces MAE by

approximately 25%, RMSE by 20%, and MAPE by 26%, demonstrating significant improvement in prediction accuracy.

The Boosting model shows better performance than ARIMA due to its ability to capture non-linear relationships. However, it still underperforms compared to the Hybrid model, which combines the strengths of both approaches.

These findings validate the effectiveness of the proposed hybrid approach for time-series forecasting in complex, real-world scenarios.

The comparative results demonstrate that the Hybrid model significantly outperforms individual models across all evaluation metrics. Specifically, the Hybrid model reduces Mean Absolute Error (MAE) by approximately 25%, Root Mean Square Error (RMSE) by 20%, and Mean Absolute Percentage Error (MAPE) by 26% when compared to the ARIMA model.

While the Boosting model performs better than ARIMA due to its ability to capture non-linear relationships, it still falls short of the Hybrid approach. The improved performance of the Hybrid model confirms that combining linear and non-linear modeling techniques leads to more accurate and reliable forecasts.

These results validate the effectiveness of the proposed system in handling complex real-world forecasting scenarios.

4.3 Graphical Analysis

Graphical analysis plays a vital role in understanding the performance of forecasting models by providing visual insights into trends, patterns, and prediction accuracy. In this project, various graphs are used to compare the outputs of different models and to evaluate the effectiveness of the proposed hybrid forecasting framework.

One of the primary visualizations used is the **Predicted vs Actual Values graph**, which compares the forecasted values generated by the models with the actual observed data. This graph helps in assessing how closely the predictions follow the real trends. It is observed that the proposed hybrid model produces predictions that align more closely with the actual values compared to the ARIMA and boosting models, indicating higher accuracy.

Another important visualization is the **Error Comparison Graph**, which illustrates the values of performance metrics such as MAE, RMSE, and MAPE for different models. This graph clearly shows that the proposed hybrid model achieves lower error values,

confirming its superior performance over traditional approaches.

The **Trend Analysis Graph** is used to examine long-term patterns in the automotive market data. This visualization highlights seasonal variations, growth trends, and fluctuations in demand. The hybrid model demonstrates a better ability to capture these patterns due to its integration of both statistical and machine learning techniques.

Additionally, a **Performance Improvement Graph** is used to compare model performance before and after optimization. This graph shows how tuning model parameters and integrating retrieval-based intelligence contribute to improved prediction accuracy.

Visualizations also include outputs from the forecasting system, such as dashboards or plotted time-series forecasts, which provide a clear representation of predicted future trends. These graphical outputs enhance the interpretability of the system and make it easier for stakeholders to understand and analyze the results.

Overall, graphical analysis confirms that the proposed framework not only improves numerical performance metrics but also provides more reliable and visually interpretable predictions. These visual insights play a crucial role in validating the effectiveness of the forecasting model and supporting data-driven decision-making.




Figure 4.3.(a): Trend Analysis of Automotive Market

Figure 4.3(a) This figure illustrates the trend of the automotive market over time, highlighting patterns such as growth, fluctuations, and seasonal variations. It provides insights into market behavior and supports the forecasting model in identifying underlying trends for accurate future predictions.




Figure 4.3(b): Performance Improvement After Optimization

Figure 4.3(b): This figure illustrates the improvement in model performance after optimization, showing a reduction in error metrics such as MAE, RMSE, and MAPE. The results indicate that the optimized and hybrid approach achieves significantly better accuracy compared to the initial models.




Figure 4.3(c): Visualization of Forecast Output Dashboard

Figure 4.3(c) This figure presents the forecast output dashboard, displaying predicted trends, historical data, and key performance indicators in a unified interface. It enables users to interpret model results visually and supports effective decision-making through interactive insights.



Figure 4.3.(d): Policy Scenario Simulation and Forecast Analysis

Figure 4.3.4 illustrates the policy simulation and scenario-based forecasting results, comparing baseline and policy-driven EV market growth. It highlights cumulative gains, growth trends, and milestone comparisons, demonstrating the impact of policy interventions on long-term market expansion.




Figure 4.3.(e): Policy Simulation Results Summary and Strategic Insights

Figure 4.3.(e) presents the summarized results of the policy simulation, highlighting the comparison between baseline and policy-driven scenarios. It includes key metrics such as projected market size, growth uplift, and additional units generated. The figure also provides strategic insights derived from the simulation, supporting data-driven decision-making.

4.4 Performance Improvement After Optimization

This section analyzes the improvement in forecasting performance after applying hybrid modeling techniques. The integration of ARIMA and Boosting models significantly reduced prediction errors compared to individual models. The hybrid approach effectively captures both linear and non-linear patterns, leading to improved accuracy and robustness.

Table 4.4: Performance Improvement After Optimization

Model	MAE	RMSE	MAPE (%)
ARIMA	200.5	220.1	7.1
Boosting Model	180.3	195.7	6.5
Hybrid Model	150.25	175.43	5.2

Table 4.4 presents the improvement in model performance after optimization. The results indicate that the proposed hybrid model achieves the highest improvement across all evaluation metrics, significantly outperforming both ARIMA and XGBoost models.

The experimental results demonstrate that the Hybrid model outperforms both ARIMA and Boosting models across all evaluation metrics. The Hybrid model reduces RMSE by approximately 20% compared to the ARIMA model, indicating improved prediction accuracy.

Furthermore, the integration of policy-based scenarios shows a significant increase in EV adoption. The policy-driven scenario results in higher projected sales compared to the baseline, highlighting the impact of government interventions on market growth.

These results validate the effectiveness of combining statistical and machine learning approaches with contextual information.

CHAPTER 5

CONCLUSION

The primary objective of this project was to develop an intelligent and efficient forecasting system for analyzing trends in the automotive market by integrating structured and unstructured data sources. The proposed **Agentic Retrieval-Augmented Forecasting Framework** successfully addresses the limitations of traditional forecasting approaches by combining statistical models, machine learning techniques, and retrieval-based intelligence.

Traditional models such as ARIMA provide a strong foundation for time-series forecasting but are limited in capturing complex non-linear relationships and external influencing factors. Machine learning models, particularly boosting algorithms, enhance prediction accuracy by modeling non-linear patterns. However, both approaches primarily rely on structured data and do not incorporate contextual knowledge from external sources.

To overcome these limitations, the proposed framework integrates **Retrieval-Augmented Generation (RAG)**, enabling the system to extract relevant insights from policy documents, industry reports, and market analyses. This integration allows the forecasting model to consider real-world factors that influence market behavior, resulting in more accurate and context-aware predictions.

Additionally, the use of an **agent-based architecture** ensures efficient coordination between different components of the system, including data preprocessing, feature engineering, document retrieval, and forecasting. This modular approach improves scalability, adaptability, and overall system performance.

The evaluation of the system using performance metrics such as **MAE, RMSE, and MAPE** demonstrates that the proposed hybrid model outperforms traditional and standalone machine learning approaches. The results confirm that combining statistical, machine learning, and retrieval-based techniques significantly improves forecasting accuracy and robustness.

In conclusion, this project presents a comprehensive and scalable solution for intelligent forecasting in the automotive domain. The integration of artificial intelligence with domain-specific knowledge provides a strong foundation for future advancements in predictive analytics and decision support systems.

5.1 Limitations

Despite the effectiveness of the proposed system, several limitations exist. The model performance is dependent on the quality and availability of input data, and the dataset used in this study is relatively limited.

The Retrieval-Augmented Generation (RAG) component is implemented using basic similarity techniques and does not incorporate large-scale vector databases, which may limit scalability.

Additionally, real-time data integration is not implemented, which may reduce adaptability to rapidly changing market conditions.

Hyperparameter tuning for machine learning models is limited, and further optimization could improve performance. Future work can address these limitations to enhance system efficiency and scalability.

CHAPTER 6

FUTURE SCOPE

The proposed Agentic Retrieval-Augmented Forecasting Framework provides a strong foundation for intelligent and context-aware prediction systems. However, there are several opportunities for further enhancement and expansion to improve its performance, scalability, and real-world applicability.

One of the key areas for future work is the integration of **advanced deep learning models**, such as Long Short-Term Memory (LSTM) networks and Transformer-based architectures. These models can capture complex temporal dependencies and long-term patterns more effectively, thereby further improving forecasting accuracy.

Another important direction is the incorporation of **real-time data processing**. Currently, the system operates on historical datasets, but integrating real-time data streams such as live market updates, economic indicators, and policy changes can make the forecasting system more dynamic and responsive to current conditions.

The framework can also be extended by enhancing the **Retrieval-Augmented Generation (RAG)** module with more sophisticated retrieval techniques and larger knowledge bases. This would enable the system to access a wider range of information sources and improve the quality of contextual insights used in forecasting.

In addition, future work can focus on developing a **user-friendly interface or web-based application** that allows stakeholders to interact with the system, visualize predictions, and analyze trends more effectively. This would make the system more accessible and practical for real-world use.

Another potential improvement is the implementation of **automated model optimization and hyperparameter tuning**, which can enhance model performance without requiring manual intervention. Techniques such as AutoML can be utilized for this purpose.

The system can also be expanded to support **multi-domain forecasting**, enabling its application in other industries such as finance, healthcare, and energy. This would demonstrate the generalizability

and scalability of the proposed approach.

Finally, future research can focus on improving the **interpretability and explainability** of the model by providing more detailed insights into the factors influencing predictions. This is particularly important for decision-making applications where transparency is essential.

In conclusion, while the proposed framework demonstrates significant improvements over traditional methods, these future enhancements can further strengthen its capabilities and broaden its applicability in intelligent forecasting systems

CHAPTER 7

REFERENCES

- [1] G. E. P. Box, G. M. Jenkins, and G. C. Reinsel, *Time Series Analysis: Forecasting and Control*, 5th ed. Hoboken, NJ, USA: Wiley, 2015.
- [2] J. Brownlee, *Introduction to Time Series Forecasting with Python*. Machine Learning Mastery, 2017.
- [3] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*, 2nd ed. New York, NY, USA: Springer, 2009.
- [4] J. H. Friedman, “Greedy function approximation: A gradient boosting machine,” *Annals of Statistics*, vol. 29, no. 5, pp. 1189–1232, 2001.
- [5] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proc. 22nd ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [6] S. Hochreiter and J. Schmidhuber, “Long short-term memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [7] A. Vaswani et al., “Attention is all you need,” in *Proc. Advances in Neural Information Processing Systems (NeurIPS)*, 2017, pp. 5998–6008.
- [8] P. Lewis et al., “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Proc. Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [9] C. D. Manning, P. Raghavan, and H. Schütze, *Introduction to Information Retrieval*. Cambridge, U.K.: Cambridge Univ. Press, 2008.
- [10] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. Cambridge, MA, USA: MIT Press, 2016.
- [11] Government of India, “Automotive Industry Report,” Ministry of Heavy Industries, New Delhi, India.
- [12] NITI Aayog, “India’s Electric Vehicle Roadmap,” Government of India.
- [13] International Energy Agency (IEA), “Global EV Outlook,” Paris, France.
- [14] McKinsey & Company, “Future of Mobility Report,” Global Industry Analysis.

APPENDIX 1

This appendix provides additional details regarding the implementation of the proposed Agentic Retrieval-Augmented Forecasting Framework, including the tools, technologies, and system workflow used in the project.

A.1 Tools and Technologies Used

The implementation of the system utilizes a combination of programming languages, libraries, and frameworks for data processing, modeling, and visualization:

- **Programming Language:** Python
 - **Data Processing Libraries:** Pandas, NumPy
 - **Visualization Tools:** Matplotlib
 - **Machine Learning Libraries:** Scikit-learn, XGBoost
 - **Time-Series Modeling:** Statsmodels (ARIMA)
 - **Natural Language Processing:** NLTK / Transformers
 - **RAG Implementation:** Vector databases and embedding models
 - **Development Environment:** Jupyter Notebook / Google Colab
-

A.2 System Architecture Overview

The system is designed using a modular architecture consisting of the following components:

1. **Data Collection Module**
 - Collects structured data (sales, economic indicators)
 - Gathers unstructured documents (policies, reports)
2. **Data Preprocessing Module**
 - Cleans and transforms structured data
 - Processes textual data using NLP techniques
3. **Feature Engineering Module**
 - Extracts time-series and contextual features
 - Generates embeddings for textual data
4. **Retrieval Module (RAG)**
 - Stores document embeddings in a vector database
 - Retrieves relevant documents based on queries
5. **Forecasting Module**
 - Implements ARIMA for linear trend modeling
 - Uses boosting models for non-linear prediction
6. **Agent-Based Coordination Module**

- Manages workflow between all components
- Ensures efficient data flow and processing

7. Output and Visualization Module

- Generates predictions and evaluation metrics
- Displays results using graphs and charts

A.3 Workflow of the System

The overall workflow of the system is as follows:

1. Data is collected from multiple sources (structured and unstructured).
2. Preprocessing is applied to clean and transform the data.
3. Features are extracted and prepared for modeling.
4. Relevant documents are retrieved using the RAG module.
5. Forecasting models (ARIMA + boosting) generate predictions.
6. Results are evaluated using performance metrics.
7. Final outputs are visualized for analysis and interpretation.

A.4 Implementation Highlights

- Hybrid modeling improves accuracy by combining statistical and ML approaches.
- RAG enhances predictions using external knowledge sources.
- Agent-based design ensures modularity and scalability.
- System is flexible and can be extended for real-time applications.

APPENDIX 2

This appendix presents sample code snippets used in the implementation of the proposed **Agentic Retrieval-Augmented Forecasting Framework**. The code demonstrates key steps including data preprocessing, ARIMA modeling, boosting model implementation, and basic retrieval operations.

A.2.1 Data Preprocessing

```
import pandas as pd
import numpy as np

# Load dataset
data = pd.read_csv("automotive_data.csv")

# Handling missing values
data.fillna(method='ffill', inplace=True)

# Convert date column to datetime
data['Date'] = pd.to_datetime(data['Date'])

# Set date as index
data.set_index('Date', inplace=True)

# Normalize data
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
data_scaled = scaler.fit_transform(data)
```

A.2.2 ARIMA Model Implementation

```
from statsmodels.tsa.arima.model import ARIMA

# Fit ARIMA model
model = ARIMA(data['Sales'], order=(2,1,2))
model_fit = model.fit()

# Forecast future values
forecast = model_fit.forecast(steps=10)
print(forecast)
```

A.2.3 Boosting Model (XGBoost)

```

from xgboost import XGBRegressor
from sklearn.model_selection import train_test_split

# Prepare features and target
X = data.drop('Sales', axis=1)
y = data['Sales']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Train model
model = XGBRegressor()
model.fit(X_train, y_train)

# Predictions
predictions = model.predict(X_test)

```

A.2.4 Performance Evaluation

```

from sklearn.metrics import mean_absolute_error, mean_squared_error

mae = mean_absolute_error(y_test, predictions)
rmse = np.sqrt(mean_squared_error(y_test, predictions))

print("MAE:", mae)
print("RMSE:", rmse)

```

A.2.5 Basic Retrieval (RAG Concept)

```

from sklearn.metrics.pairwise import cosine_similarity
from sentence_transformers import SentenceTransformer

# Sample documents
documents = ["EV policy update", "Automotive sales growth", "Fuel price impact"]

# Convert text to embeddings
model = SentenceTransformer('all-MiniLM-L6-v2')
doc_embeddings = model.encode(documents)

# Query
query = "electric vehicle policy"
query_embedding = model.encode([query])

# Similarity search
similarity = cosine_similarity(query_embedding, doc_embeddings)
print("Most relevant document:", documents[similarity.argmax()])

```

A.2.6 Visualization

```
import matplotlib.pyplot as plt

plt.plot(y_test.values, label='Actual')
plt.plot(predictions, label='Predicted')
plt.legend()
plt.title("Actual vs Predicted")
plt.show()
```

PAPER PUBLICATION STATUS



PLAGIARISM REPORT



2% detected as AI

The percentage indicates the combined amount of likely AI-generated text as well as likely AI-generated text that was also likely AI-paraphrased.

Caution: Review required.

It is essential to understand the limitations of AI detection before making decisions about a student's work. We encourage you to learn more about Turnitin's AI detection capabilities before using the tool.

Detection Groups



82 AI-generated only **2%**

Likely AI-generated text from a large-language model.



0 AI-generated text that was AI-paraphrased **0%**

Likely AI-generated text that was likely revised using an AI-paraphrase tool or word spinner.

Disclaimer

Our AI writing assessment is designed to help educators identify text that might be prepared by a generative AI tool. Our AI writing assessment may not always be accurate (i.e., our AI models may produce either false positive results or false negative results), so it should not be used as the sole basis for adverse actions against a student. It takes further scrutiny and human judgment in conjunction with an organization's application of its specific academic policies to determine whether any academic misconduct has occurred.

Frequently Asked Questions

How should I interpret Turnitin's AI writing percentage and false positives?

The percentage shown in the AI writing report is the amount of qualifying text within the submission that Turnitin's AI writing detection model determines was either likely AI-generated text from a large-language model or likely AI-generated text that was likely revised using an AI paraphrase tool or word spinner.

False positives (incorrectly flagging human-written text as AI-generated) are a possibility in AI models.

AI detection scores under 20%, which we do not surface in new reports, have a higher likelihood of false positives. To reduce the likelihood of misinterpretation, no score or highlights are attributed and are indicated with an asterisk in the report (*%).

The AI writing percentage should not be the sole basis to determine whether misconduct has occurred. The reviewer/instructor should use the percentage as a means to start a formative conversation with their student and/or use it to examine the submitted assignment in accordance with their school's policies.

What does 'qualifying text' mean?

Our model only processes qualifying text in the form of long-form writing. Long-form writing means individual sentences contained in paragraphs that make up a longer piece of written work, such as an essay, a dissertation, or an article, etc. Qualifying text that has been determined to be likely AI-generated will be highlighted in cyan in the submission, and likely AI-generated and then likely AI-paraphrased will be highlighted purple.

Non-qualifying text, such as bullet points, annotated bibliographies, etc., will not be processed and can create disparity between the submission highlights and the percentage shown.



6% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Filtered from the Report

▸ Quoted Text

Match Groups

-  **37 Not Cited or Quoted 6%**
Matches with neither in-text citation nor quotation marks
-  **0 Missing Quotations 0%**
Matches that are still very similar to source material
-  **0 Missing Citation 0%**
Matches that have quotation marks, but no in-text citation
-  **0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 4%  Internet sources
- 4%  Publications
- 2%  Submitted works (Student Papers)

Integrity Flags

0 Integrity Flags for Review

No suspicious text manipulations found.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.